

# Flash Attention: Fast and Memory-Efficient Attention

## Understanding GPU Memory Optimization

Akash Malhotra, PhD at LISN

November 11, 2025

# Outline

- 1 Background: CUDA Kernels
- 2 GPU Memory Hierarchy
- 3 Self-Attention Mechanism
- 4 Naive Attention Implementation
- 5 Flash Attention
- 6 Comparison
- 7 Results & Conclusion

# What is a CUDA Kernel?

- A **CUDA kernel** is a function that runs on the GPU
- Written in CUDA C/C++ and executed by thousands of threads in parallel
- Each thread executes the same code but on different data (SIMT model)

## Key Characteristics

- **Small programs:** Designed for specific, focused operations
- **Massively parallel:** Can launch millions of threads
- **Kernel boundaries:** After execution, control returns to CPU

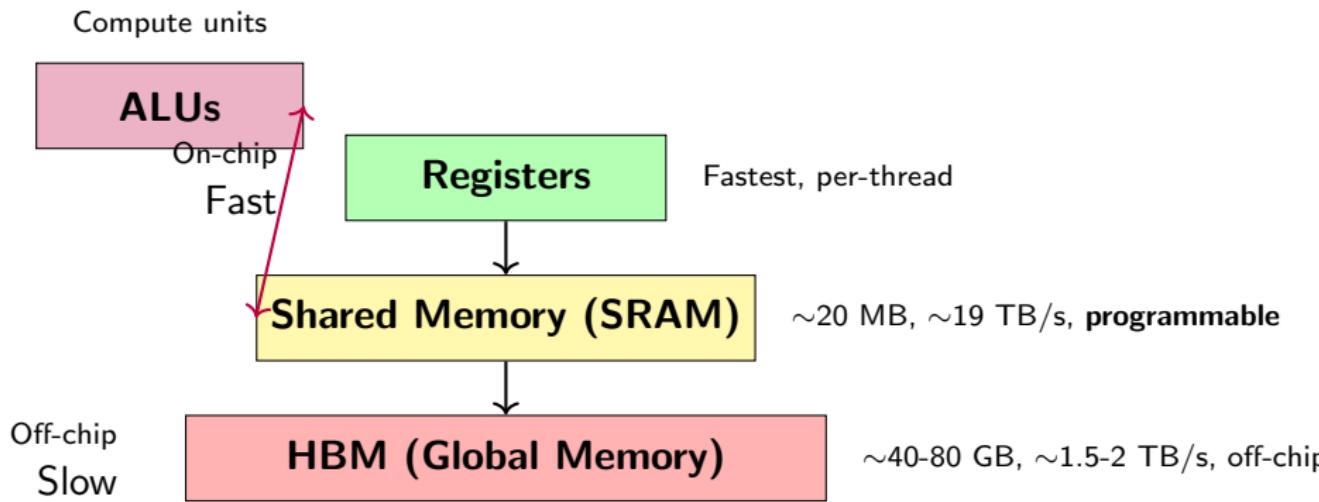
# Simple CUDA Kernel Example

```
// CUDA kernel definition
__global__ void vectorAdd(float* A, float* B, float* C, int
N) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < N) {
        C[idx] = A[idx] + B[idx]; // Each thread adds one
                                   element
    }
}

// Launch kernel from CPU (<<< >>> syntax)
vectorAdd<<<numBlocks, threadsPerBlock>>>(A, B, C, N);
```

Each thread computes one element of the result in parallel!

# GPU Memory Hierarchy



**Key Insight:** HBM is ~10x slower than SRAM! When transferring data from HBM, ALUs sit idle waiting. Flash Attention keeps data in SRAM to keep ALUs busy.

*Note: L2 cache exists but is hardware-managed and transparent to programmers*

# Why is HBM Slower?

## Physical Reasons:

- ① **Size:** Larger memory → longer access time
- ② **Location:** Off-chip → signal must travel further
- ③ **Latency:**
  - SRAM: ~nanoseconds
  - HBM: ~microseconds (1000x slower)

### The Problem

When data is transferred HBM ↔ SRAM, the ALUs (arithmetic units) sit **idle**, waiting for data.

This is called being **memory-bound**.

# Self-Attention Mechanism

**First:** Compute query, key, value matrices from input  $X$ :

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

**Then:** Compute attention using  $Q, K, V \in \mathbb{R}^{N \times d}$ :

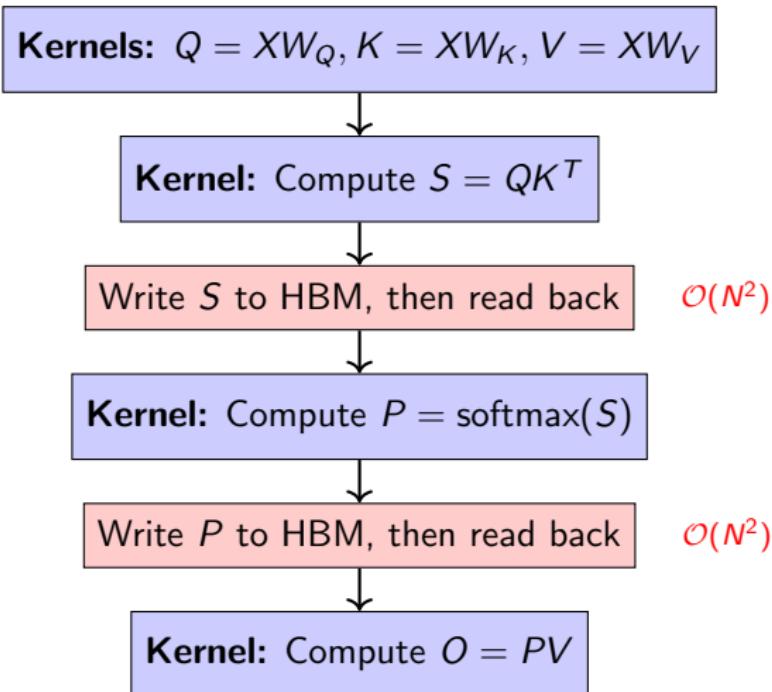
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

## Step by Step

- ① Compute  $S = QK^T \in \mathbb{R}^{N \times N}$  (attention scores)
- ② Compute  $P = \text{softmax}(S) \in \mathbb{R}^{N \times N}$  (attention weights)
- ③ Compute  $O = PV \in \mathbb{R}^{N \times d}$  (output)

**Memory Issue:** The attention matrices  $S$  and  $P$  are both  $N \times N$  — for  $N = 4096$ , this is  $\sim 67$  MB each!

# Naive CUDA Implementation: Separate Kernels



**The "Swap Memory" Problem:** HBM acts like swap space —  $S$  and  $P$  are written to HBM, SRAM flushed, then read back for next kernel!

# Flash Attention: Key Ideas

## Main Goal

**Reduce HBM access** by keeping computations in fast SRAM

- ① **Kernel Fusion:** Combine all attention operations into a single kernel
- ② **Tiling:** Divide  $Q, K, V$  into small blocks that fit in SRAM
- ③ **Recomputation:** Use online softmax algorithm to avoid storing full attention matrix
- ④ **Streaming:** Process tiles incrementally, accumulating results

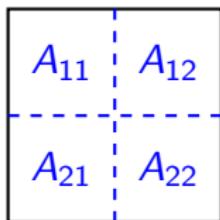
**Result:** Never materialize the full  $N \times N$  attention matrix in HBM!

# Tiling Strategy: Matrix Multiplication Example

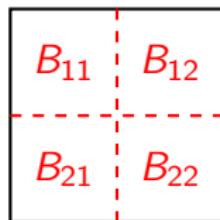
## Goal

Compute  $C = A \times B$  where  $A, B$  are  $4 \times 4$  matrices using  $2 \times 2$  tiles

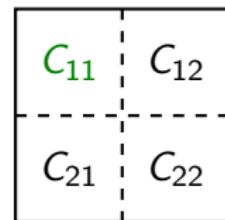
Matrix A:



Matrix B:



Result C:



Computing tile  $C_{11}$ :

$$C_{11} = A_{11} \times B_{11} + A_{12} \times B_{21}$$

**Key Point:** Each  $2 \times 2$  tile fits in SRAM. Load tiles from HBM, compute in SRAM, accumulate!

# Flash Attention Algorithm

## One Fused Kernel with Tiling

For each block  $Q_i$ :

- ① Load  $Q_i$  into SRAM
- ② Initialize output block  $O_i = 0$ , running statistics
- ③ For each block  $K_j, V_j$ :
  - Load  $K_j, V_j$  into SRAM
  - Compute  $S_{ij} = Q_i K_j^T$  (in SRAM)
  - Update online softmax statistics
  - Accumulate:  $O_i += \text{softmax}(S_{ij}) V_j$
- ④ Write final  $O_i$  to HBM

### Memory I/O:

- Read  $Q, K, V$  tiles once
- Write  $O$  once
- **No intermediate writes!**

### Key Benefits:

- Everything stays in SRAM
- ALUs stay busy
- Same  $\mathcal{O}(N^2d)$  FLOPs

# Naive vs Flash Attention: Side by Side

Aspect	Naive	Flash Attention
# of Kernels	3 separate	1 fused
HBM Reads	$\mathcal{O}(N^2)$	$\mathcal{O}(Nd)$
HBM Writes	$\mathcal{O}(N^2)$	$\mathcal{O}(Nd)$
Store full $S, P$ ?	Yes	No
SRAM Usage	Flushed between kernels	Persistent
Bottleneck	Memory-bound	Compute-bound
Speedup	1x (baseline)	2-4x faster

## The Key Difference

Flash Attention does the **same amount of computation** but with **much less memory traffic!**

# The Online Softmax Trick

**Challenge:** Standard softmax requires all values before computing:

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

**Solution:** Online/incremental softmax algorithm

- Compute softmax **incrementally** as new blocks arrive
- Maintain running max and sum of exponentials
- Update previous results when processing new blocks

This enables computing attention without storing the full  $N \times N$  matrix!

*Details: "FlashAttention: Fast and Memory-Efficient Exact Attention" (Dao et al., 2022)*

# Results & Key Takeaways

## Performance:

- **2-4x faster** than standard attention
- Greater speedup for longer sequences
- Memory:  $\mathcal{O}(Nd)$  vs  $\mathcal{O}(N^2)$
- Train with 4-8x longer contexts

**Impact:** Used in GPT-4, Claude, LLaMA, etc.

## Key Lessons:

- ➊ Memory hierarchy matters (SRAM vs HBM)
- ➋ Memory-bound vs compute-bound
- ➌ Kernel fusion + tiling keeps data in fast memory
- ➍ Algorithmic innovation (online softmax)
- ➎ Same math, better implementation

# Further Reading

- **Original Paper:** "FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness"  
Dao, Fu, Ermon, Rudra, Ré (2022)  
<https://arxiv.org/abs/2205.14135>
- **FlashAttention-2:** Further optimizations  
Dao (2023)  
<https://arxiv.org/abs/2307.08691>
- **CUDA Programming Guide:**  
<https://docs.nvidia.com/cuda/>

**Questions?**