

Solutions to Transformer Architecture Exercises

Solution Manual

November 16, 2025

Contents

1 Exercise 1: Upper Bound on Attention Coefficients	2
2 Exercise 2: Softmax Satisfies Constraints	2
3 Exercise 3: Orthogonal Input Vectors	2
4 Exercise 4: Expected Value of Dot Product Squared	4
5 Exercise 5: Multi-Head Attention Reformulation	4
6 Exercise 6: Self-Attention as Parameter-Efficient Network	5
7 Exercise 7: Equivariance Without Positional Encoding	6
8 Exercise 8: High-Dimensional Orthogonality	7
9 Exercise 9: Linear Transformation of Concatenated Vectors	8
10 Exercise 10: Sinusoidal Encoding Properties	8

1 Exercise 1: Upper Bound on Attention Coefficients

Problem

Given $a_{nm} \geq 0$ and $\sum_m a_{nm} = 1$, show that $a_{nm} \leq 1$.

Approach

Since non-negative numbers sum to 1, no single number can exceed 1 (like sharing a pizza—no one gets more than the whole pizza).

Solution

For any coefficient a_{nk} :

$$a_{nk} + \sum_{m \neq k} a_{nm} = 1 \quad (1)$$

Since all terms are non-negative: $\sum_{m \neq k} a_{nm} \geq 0$, therefore:

$$a_{nk} = 1 - \sum_{m \neq k} a_{nm} \leq 1 \quad \checkmark \quad (2)$$

2 Exercise 2: Softmax Satisfies Constraints

Problem

Verify that $a_{nm} = \frac{\exp(\mathbf{x}_n^T \mathbf{x}_m)}{\sum_{m'} \exp(\mathbf{x}_n^T \mathbf{x}_{m'})}$ satisfies: (1) $a_{nm} \geq 0$ and (2) $\sum_m a_{nm} = 1$.

Approach

Exponential is always positive (ensures non-negativity), and we normalize by the sum (ensures sum to 1).

Solution

Non-negativity: Since $\exp(z) > 0$ for all z , both numerator and denominator are positive, so $a_{nm} > 0$. \checkmark

Normalization: The denominator is constant for fixed n :

$$\sum_m a_{nm} = \frac{\sum_m \exp(\mathbf{x}_n^T \mathbf{x}_m)}{\sum_{m'} \exp(\mathbf{x}_n^T \mathbf{x}_{m'})} = 1 \quad \checkmark \quad (3)$$

3 Exercise 3: Orthogonal Input Vectors

Problem

For orthogonal vectors ($\mathbf{x}_n^T \mathbf{x}_m = 0$ for $n \neq m$), show that self-attention returns the input unchanged: $\mathbf{Y} = \mathbf{X}$.

Approach

For orthogonal vectors, $\mathbf{X}\mathbf{X}^T$ is diagonal, so each token's query only has high similarity with itself. But softmax still spreads some probability mass to other tokens unless the diagonal values are very large.

Solution

Self-attention: $\mathbf{Y} = \text{Softmax}[\mathbf{X}\mathbf{X}^T]\mathbf{X}$

For orthogonal vectors:

$$(\mathbf{X}\mathbf{X}^T)_{nm} = \begin{cases} \|\mathbf{x}_n\|^2 & \text{if } n = m \text{ (matching token)} \\ 0 & \text{if } n \neq m \text{ (all other tokens)} \end{cases} \quad (4)$$

So $\mathbf{X}\mathbf{X}^T$ is diagonal. Applying softmax to row n :

$$\text{Row } n : [\|\mathbf{x}_n\|^2, 0, 0, \dots, 0] \quad (5)$$

After softmax:

$$a_{nn} = \frac{e^{\|\mathbf{x}_n\|^2}}{e^{\|\mathbf{x}_n\|^2} + (N - 1)}, \quad a_{nm} = \frac{1}{e^{\|\mathbf{x}_n\|^2} + (N - 1)} \text{ for } m \neq n \quad (6)$$

Key insight: Softmax always normalizes to sum to 1, so even when other entries are 0, they contribute $\exp(0) = 1$ each!

Concrete example: For $\|\mathbf{x}_n\|^2 = 1$ and $N = 10$ tokens:

$$a_{nn} = \frac{e^1}{e^1 + 9} = \frac{2.718}{11.718} \approx 0.232 \quad (7)$$

$$a_{nm} = \frac{1}{11.718} \approx 0.085 \text{ for each of the 9 other tokens} \quad (8)$$

Check: $0.232 + 9 \times 0.085 = 0.232 + 0.765 \approx 1$

So the output is:

$$\mathbf{y}_n = 0.232 \mathbf{x}_n + 0.085 \mathbf{x}_1 + \dots + 0.085 \mathbf{x}_{n-1} + 0.085 \mathbf{x}_{n+1} + \dots \quad (9)$$

This is **not** equal to \mathbf{x}_n —it's a weighted average where token n gets 23% weight and the other 9 tokens share 77%.

When does $\mathbf{Y} = \mathbf{X}$ actually hold?

Only when diagonal entries are much larger. For example, if $\|\mathbf{x}_n\|^2 = 10$:

$$a_{nn} = \frac{e^{10}}{e^{10} + 9} = \frac{22,026}{22,035} \approx 0.9996 \quad (\text{almost } 1!) \quad (10)$$

In the limit $\|\mathbf{x}_n\|^2 \rightarrow \infty$, we get $a_{nn} \rightarrow 1$ and $a_{nm} \rightarrow 0$, giving $\mathbf{Y} = \mathbf{X}$.

Intuition: Think of softmax as a "soft" version of picking the maximum. When you have scores [5, 0, 0, 0], softmax heavily favors the first entry but still gives small probabilities to others. Only when the gap is huge (like [100, 0, 0, 0]) does it become nearly one-hot.

For this exercise: The statement " $\mathbf{Y} = \mathbf{X}$ " is true *in the limit* or requires assuming unit vectors are implicitly scaled. For typical unit vectors, we get $\mathbf{Y} \approx \mathbf{X}$ with some mixing. ✓

4 Exercise 4: Expected Value of Dot Product Squared

Problem

For $\mathbf{a}, \mathbf{b} \sim \mathcal{N}(0, \mathbf{I}_D)$ independent, show $E[(\mathbf{a}^T \mathbf{b})^2] = D$.

Approach

Expand the square, use independence. Only diagonal terms survive in the double sum.

Solution

Each component: $a_i, b_i \sim \mathcal{N}(0, 1)$ with $E[a_i] = 0, E[a_i^2] = 1$.

Mean: $E[\mathbf{a}^T \mathbf{b}] = \sum_i E[a_i]E[b_i] = 0$

Second moment:

$$(\mathbf{a}^T \mathbf{b})^2 = \left(\sum_i a_i b_i \right)^2 = \sum_{i,j} a_i a_j b_i b_j \quad (11)$$

Taking expectations and using independence:

$$E[(\mathbf{a}^T \mathbf{b})^2] = \sum_{i,j} E[a_i a_j] E[b_i b_j] \quad (12)$$

Since $E[a_i a_j] = \delta_{ij}$ (equals 1 if $i = j$, else 0), only diagonal terms ($i = j$) survive:

$$E[(\mathbf{a}^T \mathbf{b})^2] = \sum_{i=1}^D 1 \cdot 1 = D \quad \checkmark \quad (13)$$

Implication: $\text{Var}(\mathbf{a}^T \mathbf{b}) = D$, so $\text{SD} = \sqrt{D}$. For $D = 512$, $\text{SD} \approx 23$ —huge! This motivates the $1/\sqrt{D_k}$ scaling in attention to keep variance at 1.

5 Exercise 5: Multi-Head Attention Reformulation

Problem

Show $\mathbf{Y} = \text{Concat}[\mathbf{H}_1, \dots, \mathbf{H}_H] \mathbf{W}^{(o)} = \sum_{h=1}^H \mathbf{H}_h \mathbf{W}_h^{(o)}$.

Approach

Block matrix multiplication: $[A|B] \begin{bmatrix} C \\ D \end{bmatrix} = AC + BD$.

Solution

Partition $\mathbf{W}^{(o)}$ vertically into H blocks: $\mathbf{W}^{(o)} = \begin{bmatrix} \mathbf{W}_1^{(o)} \\ \vdots \\ \mathbf{W}_H^{(o)} \end{bmatrix}$

Then:

$$[\mathbf{H}_1 | \cdots | \mathbf{H}_H] \begin{bmatrix} \mathbf{W}_1^{(o)} \\ \vdots \\ \mathbf{W}_H^{(o)} \end{bmatrix} = \sum_{h=1}^H \mathbf{H}_h \mathbf{W}_h^{(o)} \quad \checkmark \quad (14)$$

This is standard block matrix multiplication.

6 Exercise 6: Self-Attention as Parameter-Efficient Network Problem

Express self-attention as a fully-connected network and explain its parameter structure.

Approach

Self-attention connects every input token to every output token (fully connected), but achieves this with far fewer parameters than a standard fully-connected layer. The key is that the connection weights are computed from the input rather than being separate learned parameters.

Solution

Standard fully-connected baseline:

To map all N input tokens (total size ND) to all N output tokens, you'd need a matrix $\mathbf{W}_{\text{full}} \in \mathbb{R}^{ND \times ND}$ with N^2D^2 independent parameters.

Self-attention approach:

Self-attention computes $\mathbf{Y} = \mathbf{AV}$ where:

- $\mathbf{V} = \mathbf{X}\mathbf{W}^{(v)}$ (value transformation)
- $\mathbf{A} = \text{Softmax}[\mathbf{Q}\mathbf{K}^T]$ (attention matrix, computed from input)

We can write this as a block matrix operating on all tokens:

$$\mathbf{Y} = \begin{bmatrix} a_{11}\mathbf{W}^{(v)} & a_{12}\mathbf{W}^{(v)} & \dots & a_{1N}\mathbf{W}^{(v)} \\ a_{21}\mathbf{W}^{(v)} & a_{22}\mathbf{W}^{(v)} & \dots & a_{2N}\mathbf{W}^{(v)} \\ \vdots & \vdots & \ddots & \vdots \\ a_{N1}\mathbf{W}^{(v)} & a_{N2}\mathbf{W}^{(v)} & \dots & a_{NN}\mathbf{W}^{(v)} \end{bmatrix} \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \vdots \\ \mathbf{x}_N \end{bmatrix} \quad (15)$$

This is an $ND \times ND$ matrix—fully connected, not sparse!

Key insight—where the efficiency comes from:

1. **Parameter sharing:** Every block uses the *same* matrix $\mathbf{W}^{(v)}$, just scaled by different a_{nm} values. That's only D^2 parameters for $\mathbf{W}^{(v)}$, not N^2D^2 !
2. **Input-dependent routing:** The scalars a_{nm} are *computed from the input* using $\mathbf{Q}\mathbf{K}^T = \mathbf{X}\mathbf{W}^{(q)}\mathbf{W}^{(k)T}\mathbf{X}^T$. This adds $2D^2$ parameters for $\mathbf{W}^{(q)}$ and $\mathbf{W}^{(k)}$, but is shared across all positions.
3. **Total learned parameters:** $3D^2$ (query + key + value matrices) instead of N^2D^2 .

Analogy: Imagine a switchboard where:

- **Fully-connected:** Each pair of phones has its own dedicated wire (needs N^2 wires)
- **Self-attention:** All pairs share the same wire, but each call has a volume knob (the a_{nm}) that's automatically adjusted based on who's calling whom. You only need one wire!

Concrete numbers: For $N = 100$ tokens, $D = 512$ dimensions:

Approach	Parameters
Fully-connected	$N^2D^2 = 2,621,440,000$ (2.6 billion)
Self-attention	$3D^2 = 786,432$
Reduction	3330× fewer parameters!

Summary: Self-attention achieves full connectivity between all tokens with a *factorized, parameter-efficient* structure. The transformation is dense (all-to-all connections), but the parameterization is low-rank (shared weights + input-dependent routing). ✓

7 Exercise 7: Equivariance Without Positional Encoding

Problem

Show that without positional encoding, permuting inputs permutes outputs the same way.

Approach

Show that permuting input rows causes \mathbf{Q} , \mathbf{K} , \mathbf{V} to be permuted identically, leading to permuted outputs.

Solution

Let \mathbf{X}^π be \mathbf{X} with rows permuted by π : $(\mathbf{X}^\pi)_n = \mathbf{x}_{\pi(n)}$.

Since $\mathbf{Q} = \mathbf{X}\mathbf{W}^{(q)}$ operates row-wise:

$$(\mathbf{Q}^\pi)_n = \mathbf{x}_{\pi(n)}\mathbf{W}^{(q)} = (\mathbf{Q})_{\pi(n)} \quad (16)$$

Same for \mathbf{K}^π , \mathbf{V}^π . Therefore attention weights: $a_{nm}^\pi = a_{\pi(n),\pi(m)}$

The output:

$$(\mathbf{Y}^\pi)_n = \sum_m a_{nm}^\pi (\mathbf{V}^\pi)_m = \sum_m a_{\pi(n),\pi(m)} (\mathbf{V})_{\pi(m)} = \sum_\ell a_{\pi(n),\ell} (\mathbf{V})_\ell = (\mathbf{Y})_{\pi(n)} \quad (17)$$

Therefore $(\mathbf{Y}^\pi)_n = (\mathbf{Y})_{\pi(n)}$. ✓

Example: "The dog chased the cat" vs "cat the chased dog The" would be processed identically (just reordered). This is why positional encoding is essential!

8 Exercise 8: High-Dimensional Orthogonality

Problem

Show random unit vectors in high dimensions are nearly orthogonal.

Approach

The dot product has mean 0 and variance $1/D$, so it concentrates near 0 as $D \rightarrow \infty$.

Solution

For random unit vectors \mathbf{a}, \mathbf{b} in \mathbb{R}^D , express in orthonormal basis:

$$\mathbf{a} = \sum_i \alpha_i \mathbf{u}_i, \quad \mathbf{b} = \sum_i \beta_i \mathbf{u}_i \quad (18)$$

with $\sum_i \alpha_i^2 = \sum_i \beta_i^2 = 1$.

For isotropic distribution: $E[\alpha_i] = 0$, $E[\alpha_i^2] = 1/D$.

Dot product:

$$\mathbf{a}^T \mathbf{b} = \sum_i \alpha_i \beta_i \quad (19)$$

Mean: $E[\mathbf{a}^T \mathbf{b}] = \sum_i E[\alpha_i] E[\beta_i] = 0$

Variance: Using independence and that only $i = j$ terms survive:

$$\text{Var}(\mathbf{a}^T \mathbf{b}) = \sum_i E[\alpha_i^2] E[\beta_i^2] = D \cdot \frac{1}{D^2} = \frac{1}{D} \quad (20)$$

Therefore $\text{SD}(\cos \theta) = 1/\sqrt{D} \rightarrow 0$ as $D \rightarrow \infty$. ✓

Examples:

D	SD	Typical angle
10	0.316	72
512	0.044	87.5

In high dimensions, random vectors are almost always perpendicular.

9 Exercise 9: Linear Transformation of Concatenated Vectors

Problem

Show that $\mathbf{W}[\mathbf{x}; \mathbf{e}] = \mathbf{W}_x \mathbf{x} + \mathbf{W}_e \mathbf{e}$.

Approach

Block matrix multiplication: partition $\mathbf{W} = [\mathbf{W}_x | \mathbf{W}_e]$ to match the concatenation.

Solution

$$\mathbf{W} \begin{bmatrix} \mathbf{x} \\ \mathbf{e} \end{bmatrix} = [\mathbf{W}_x | \mathbf{W}_e] \begin{bmatrix} \mathbf{x} \\ \mathbf{e} \end{bmatrix} = \mathbf{W}_x \mathbf{x} + \mathbf{W}_e \mathbf{e} \quad \checkmark \quad (21)$$

Implication: Concatenation + linear layer \approx addition + linear layer. Transformers use addition because it's simpler (no dimension increase, fewer parameters).

10 Exercise 10: Sinusoidal Encoding Properties

Problem

Show that $\mathbf{r}_{n+k} = \mathbf{T}_k \mathbf{r}_n$ where \mathbf{T}_k depends only on k , not n .

Approach

Use trig identities: $\sin(A + B) = \sin A \cos B + \cos A \sin B$ and $\cos(A + B) = \cos A \cos B - \sin A \sin B$.

Solution

Let $\omega_i = 1/L^{i/D}$. The encoding pairs sine and cosine at each frequency.

For even i (sine):

$$r_{n+k,i} = \sin((n+k)\omega_i) = \sin(n\omega_i) \cos(k\omega_i) + \cos(n\omega_i) \sin(k\omega_i) \quad (22)$$

$$= r_{n,i} \cos(k\omega_i) + r_{n,i+1} \sin(k\omega_i) \quad (23)$$

For odd i (cosine):

$$r_{n+k,i} = \cos((n+k)\omega_{i-1}) = \cos(n\omega_{i-1}) \cos(k\omega_{i-1}) - \sin(n\omega_{i-1}) \sin(k\omega_{i-1}) \quad (24)$$

$$= r_{n,i} \cos(k\omega_{i-1}) - r_{n,i-1} \sin(k\omega_{i-1}) \quad (25)$$

Matrix form:

$$\begin{bmatrix} r_{n+k,i} \\ r_{n+k,i+1} \end{bmatrix} = \begin{bmatrix} \cos(k\omega) & \sin(k\omega) \\ -\sin(k\omega) & \cos(k\omega) \end{bmatrix} \begin{bmatrix} r_{n,i} \\ r_{n,i+1} \end{bmatrix} \quad (26)$$

This is a rotation matrix! The angle $k\omega$ depends only on k , not n .

The full transformation $\mathbf{r}_{n+k} = \mathbf{T}_k \mathbf{r}_n$ is block-diagonal with these rotation matrices. ✓

Why sine-only fails: $\sin((n+k)\omega) = \sin(n\omega) \cos(k\omega) + \cos(n\omega) \sin(k\omega)$ requires $\cos(n\omega)$, which isn't available in sine-only encoding.

Why this matters: The network can learn to attend to relative positions (e.g., "2 tokens back") uniformly across the sequence through linear transformations.